

# Horizontal Partitioning

---



**Leonard Lobel**

CTO, SLEEK TECHNOLOGIES

[lennilobel.wordpress.com](http://lennilobel.wordpress.com)



# Achieving Elastic Scale

## What Is Partitioning?

Massive scale-out within a container

## Unlimited Containers

Logical resource  
composed of multiple partitions

## Partitions

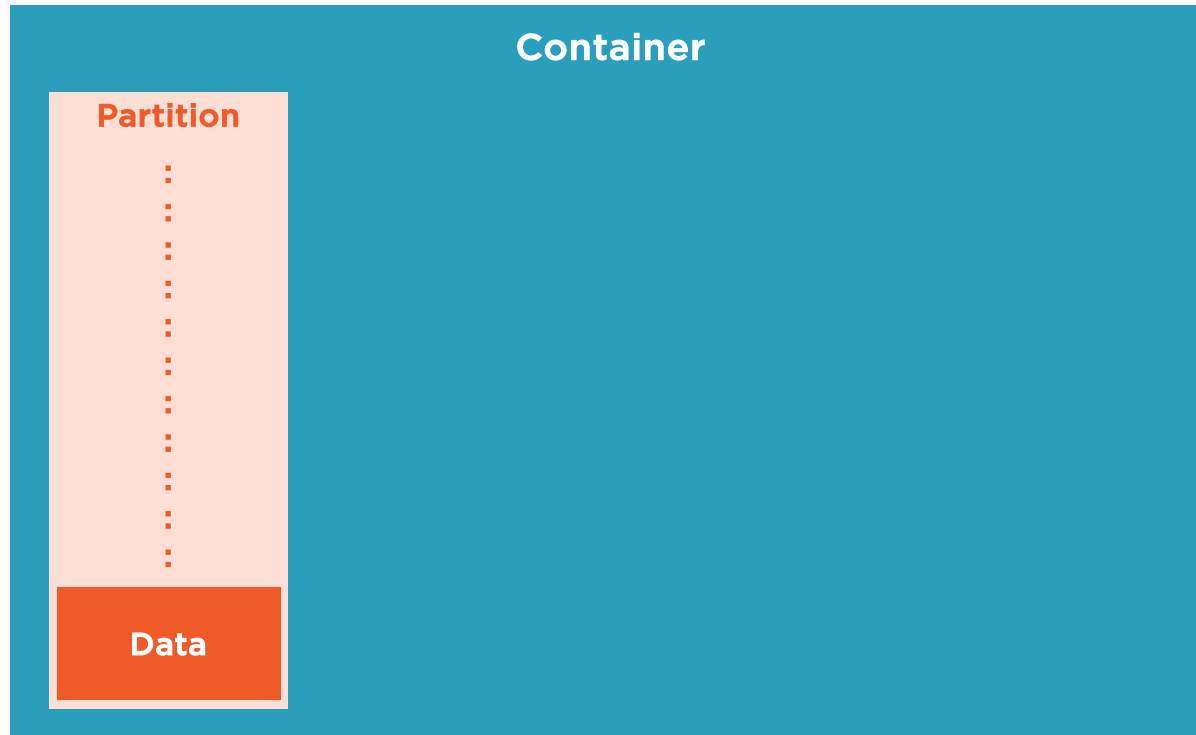
Physical fixed-capacity data buckets

## Automated Scale-out


Cosmos DB transparently splits  
partitions to manage growth



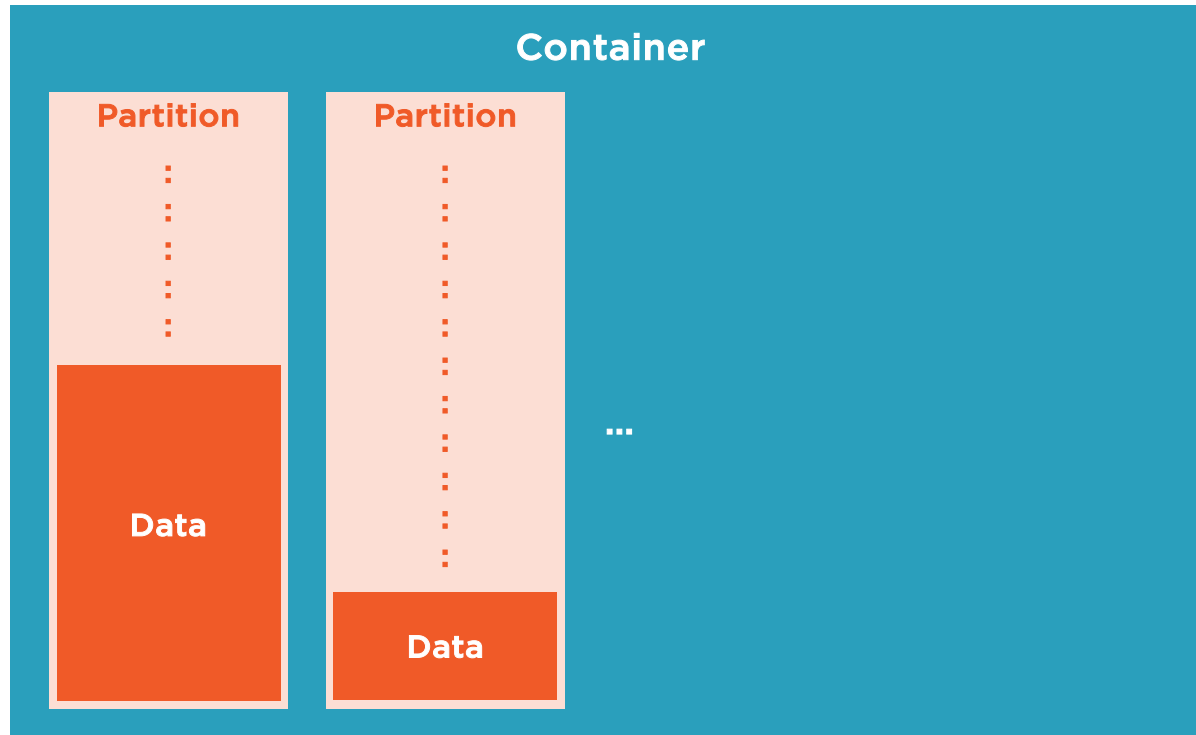
# Achieving Elastic Scale



# Achieving Elastic Scale



# Achieving Elastic Scale



# Selecting a Partition Key

## Choosing the best partition key

The right choice will deliver massive scale

## Partition key values are hashed

Hashed value determines the physical partition for storing each item

## Partitions host multiple partition keys

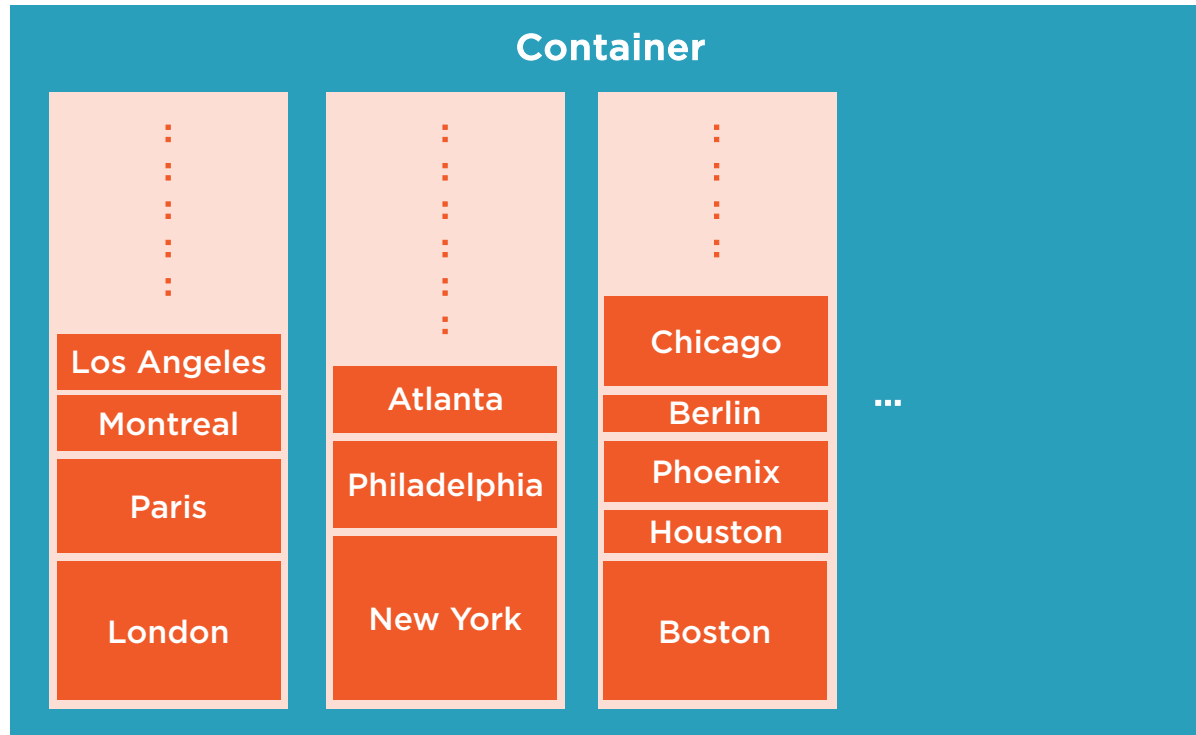
Items with the same partition key value are physically stored together on the same partition

## Two primary considerations

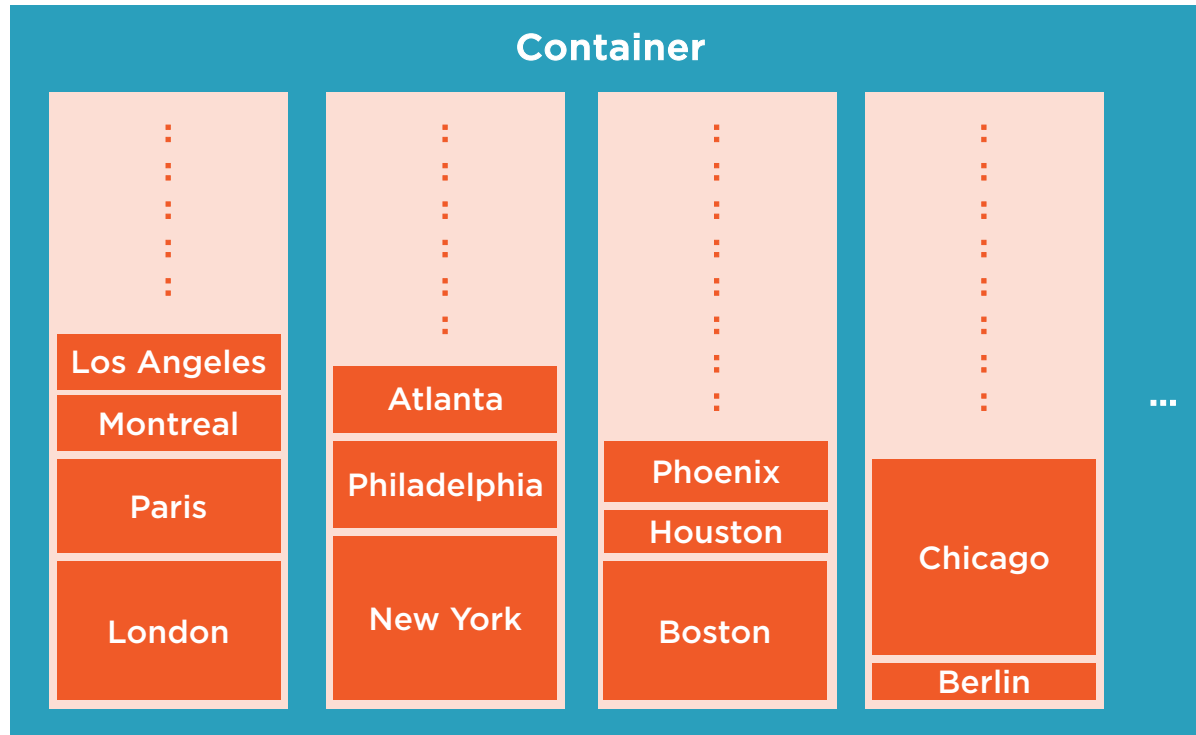
- 1) Boundary for query and transactions
- 2) No storage or performance bottlenecks



# Selecting a Partition Key



# Selecting a Partition Key





# Choosing the Right Partition Key

## Driven by data access patterns

Choose a property that groups commonly queried/updated items together

**User Profile  
Data**

User ID

**IoT  
(e.g., device state)**

Device ID

**Multi-tenant  
Architecture**

Tenant ID

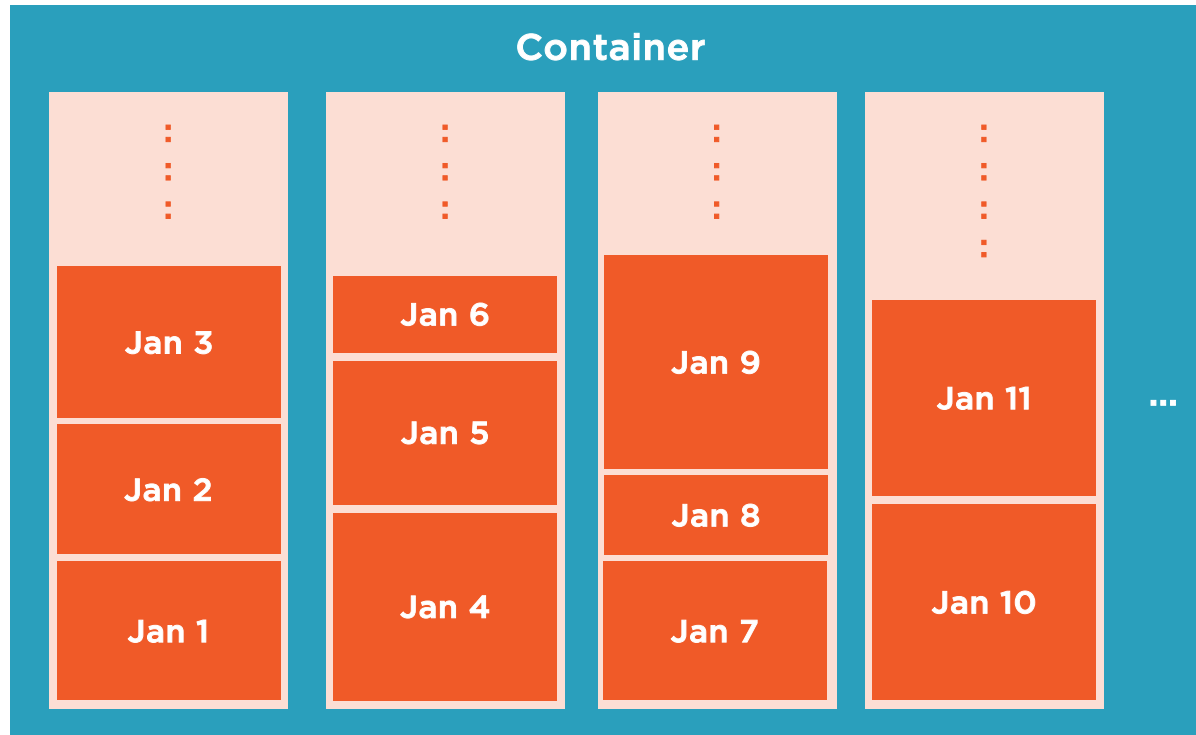


# Choosing the Right Partition Key

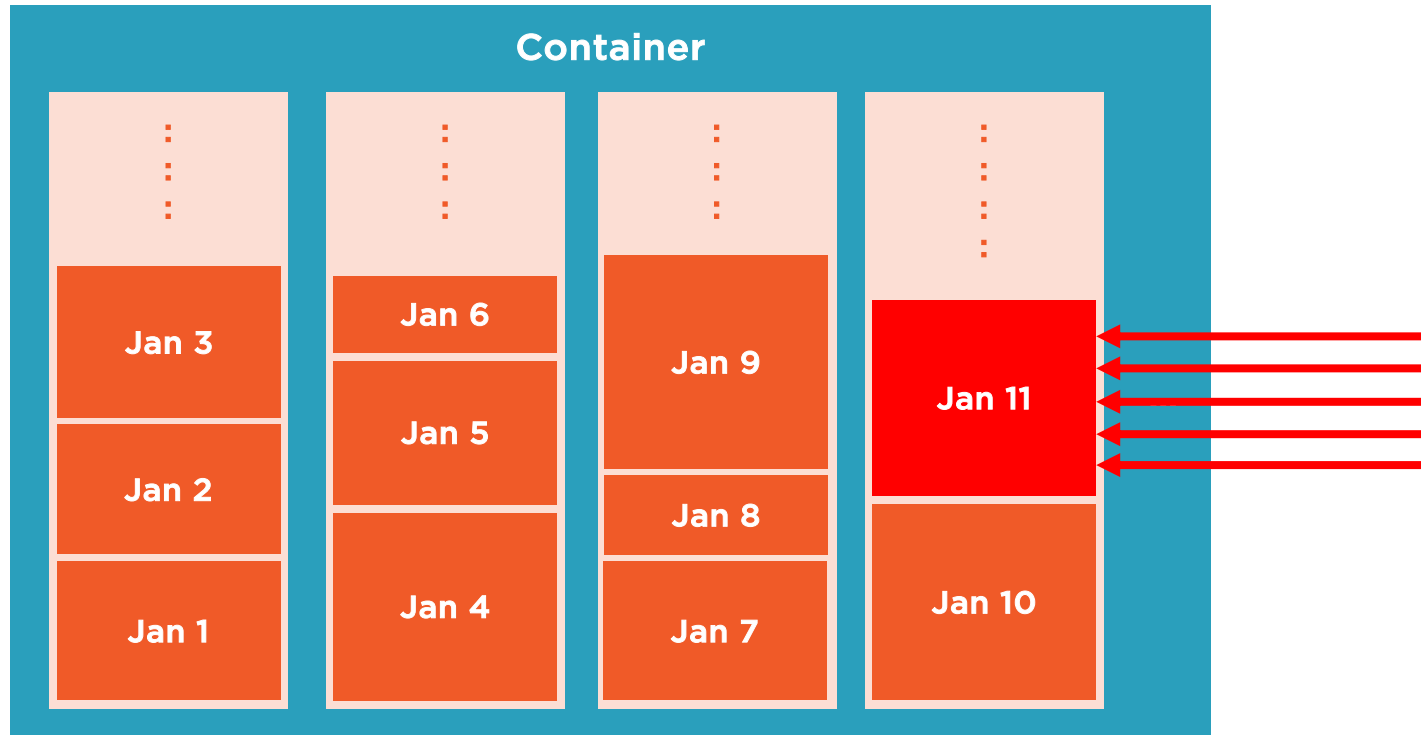
- Generally, writes should be distributed uniformly across partitions
- For example, user profile data with a user ID and creation date
  - Partitioning by creation date
    - Bad idea! All writes of the day are directed to the same partition



# Choosing the Right Partition Key



# Choosing the Right Partition Key

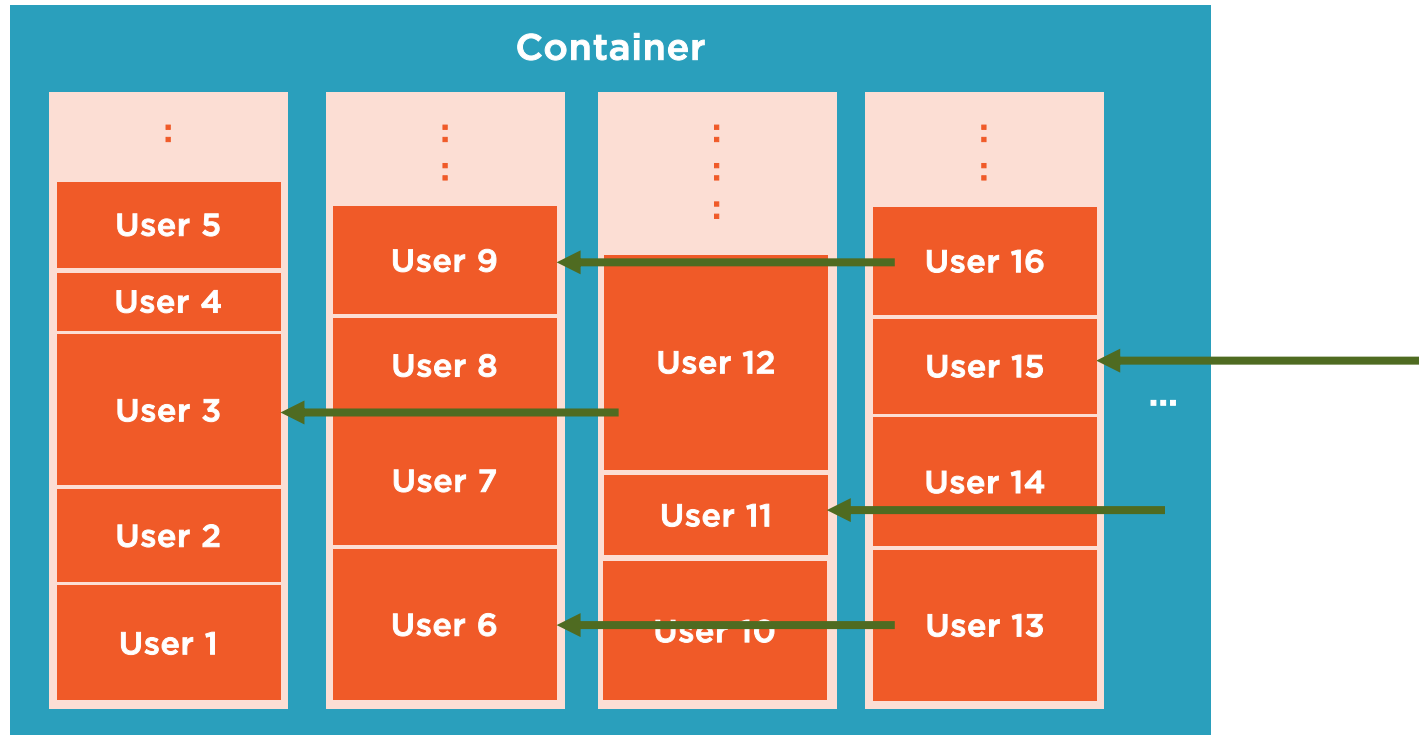


# Choosing the Right Partition Key

- Generally, writes should be distributed uniformly across partitions
- For example, user profile data with a user ID and creation date
  - Partitioning by creation date
    - Bad idea! All writes of the day are directed to the same partition
  - Partition by user ID
    - Much better! Writes are directed to different partitions per user



# Choosing the Right Partition Key

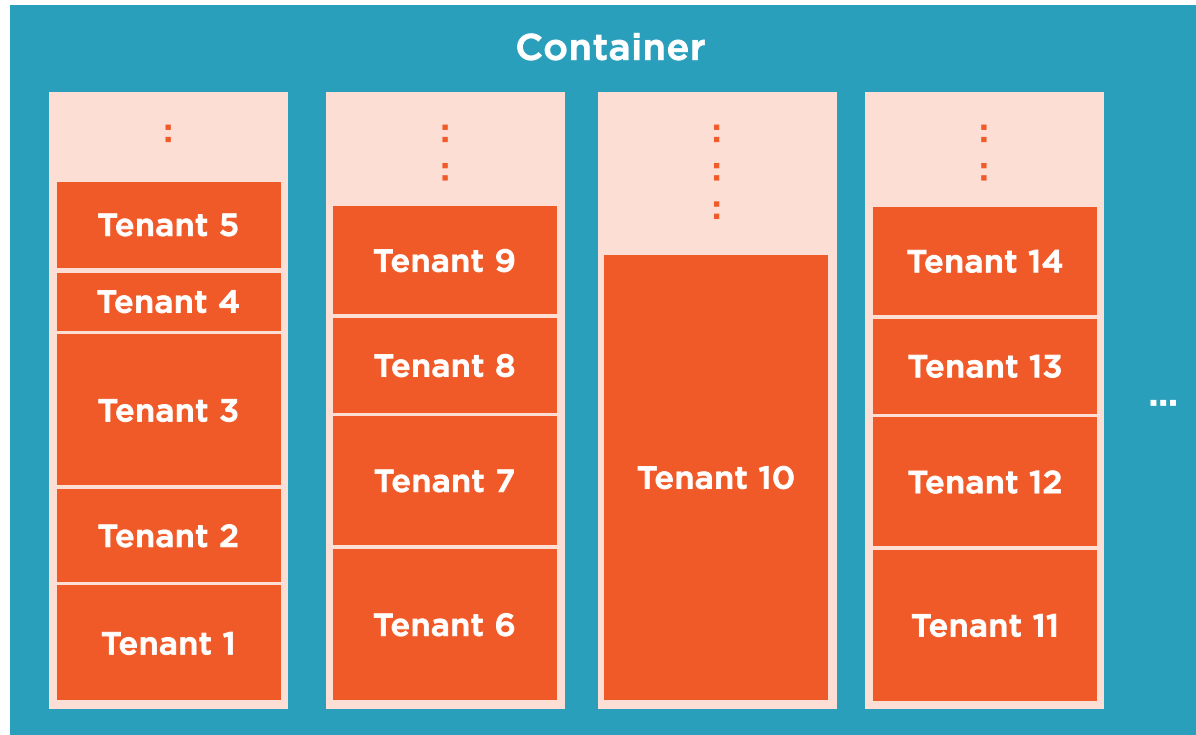


# Choosing the Right Partition Key

- Generally, writes should be distributed uniformly across partitions
- For example, user profile data with a user ID and creation date
  - Partitioning by creation date
    - Bad idea! All writes of the day are directed to the same partition
  - Partition by user ID
    - Much better! Writes are directed to different partitions per user
- Create multiple containers for varying throughput needs
  - Throughput is purchased at the container level

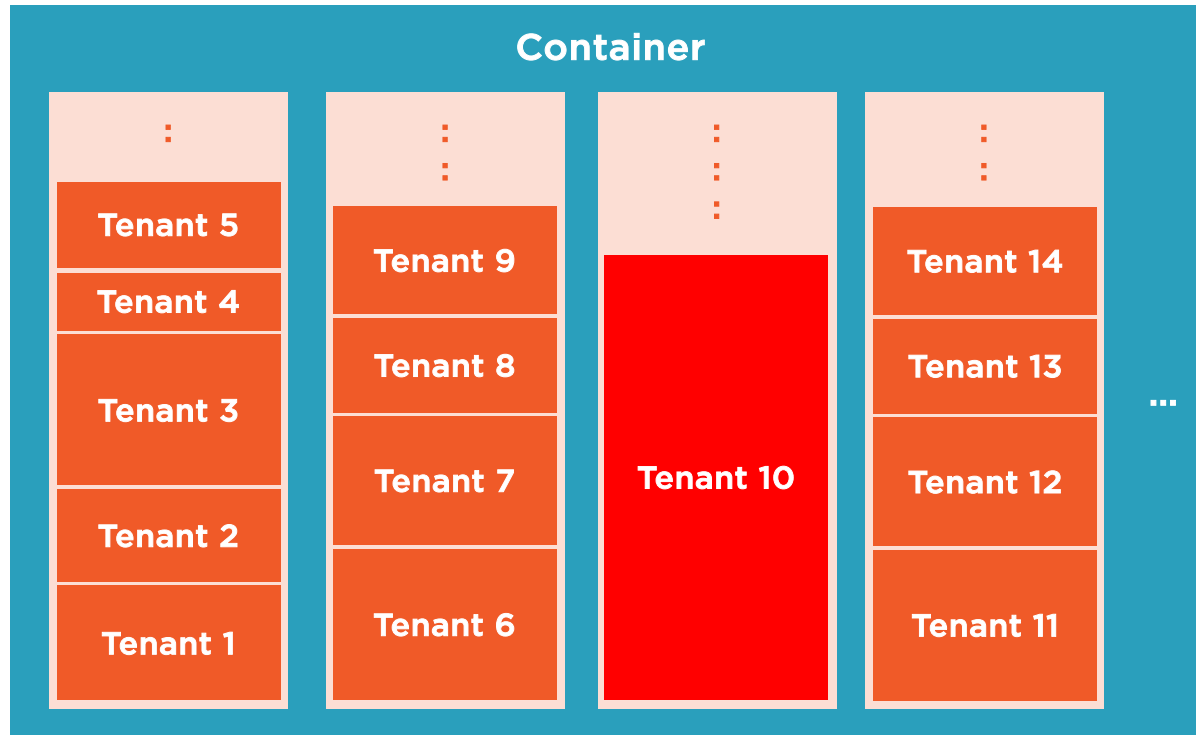


# Choosing the Right Partition Key


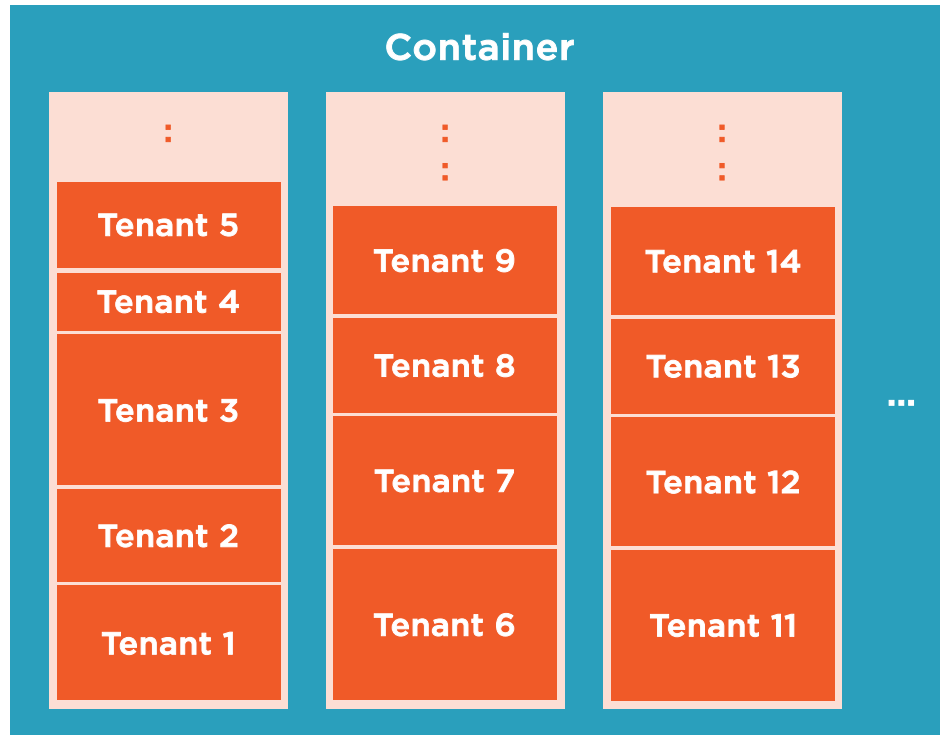




# Choosing the Right Partition Key



# Choosing the Right Partition Key



# Cross Partition Queries

## Stored Procedures

Always scoped to a  
single partition key

## Queries

Typically scoped to a  
single partition key

## Cross-Partition Queries

Span multiple  
partition keys  
Fan-out execution



```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}
```

```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList(); ≤ 7ms elapsed

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}
```

```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'"; ≤ 261ms elapsed
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}
```

```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message); ≤ 78ms elapsed
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}
```

```

using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }
}

```

≤ 78ms elapsed

ex Count = 1

<ul style="list-style-type: none"> <li>Data</li> <li>HResult</li> <li>HelpLink</li> <li>InnerException</li> <li>InnerExceptions</li> <li>Message</li> </ul>	<ul style="list-style-type: none"> <li>{System.Collections.ListDictionaryInternal}</li> <li>-2146233088</li> <li>null</li> <li>{"Cross partition query is required but disabled. Please set x-ms-documentdb-query-enablecross"</li> <li>Count = 1</li> <li>"One or more errors occurred."</li> </ul>
---	--



```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message); ≤ 78ms elapsed
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}
```

```
using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList(); ≤ 3ms elapsed

    await client.DeleteDatabaseAsync(dbUri);
}
```

```

using (var client = new DocumentClient(new Uri(endpoint), masterKey))
{
    var db = new Database { Id = "families" };
    await client.CreateDatabaseAsync(db);

    var coll = new DocumentCollection { Id = "families" };
    coll.PartitionKey.Paths.Add("/address/zipCode");

    await client.CreateDocumentCollectionAsync(dbUri, coll, new RequestOptions { OfferThroughput = 20000 });

    var sql = "SELECT * FROM c WHERE c.address.zipCode = '60603'";
    var query = client.CreateDocumentQuery<Document>(collUri, sql);
    var result = query.ToList();

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    query = client.CreateDocumentQuery<Document>(collUri, sql);
    try
    {
        result = query.ToList();
    }
    catch (Exception ex)
    {
        Console.WriteLine(ex.Message);
    }

    sql = "SELECT * FROM c WHERE c.address.city = 'Chicago'";
    var options = new FeedOptions { EnableCrossPartitionQuery = true, MaxDegreeOfParallelism = -1 };
    query = client.CreateDocumentQuery<Document>(collUri, sql, options);
    result = query.ToList();

    await client.DeleteDatabaseAsync(dbUri);
}

```

≤ 936ms elapsed

# Summary



Achieving elastic scale

Horizontal partitioning

Choosing the right partition key

Cross partition queries

